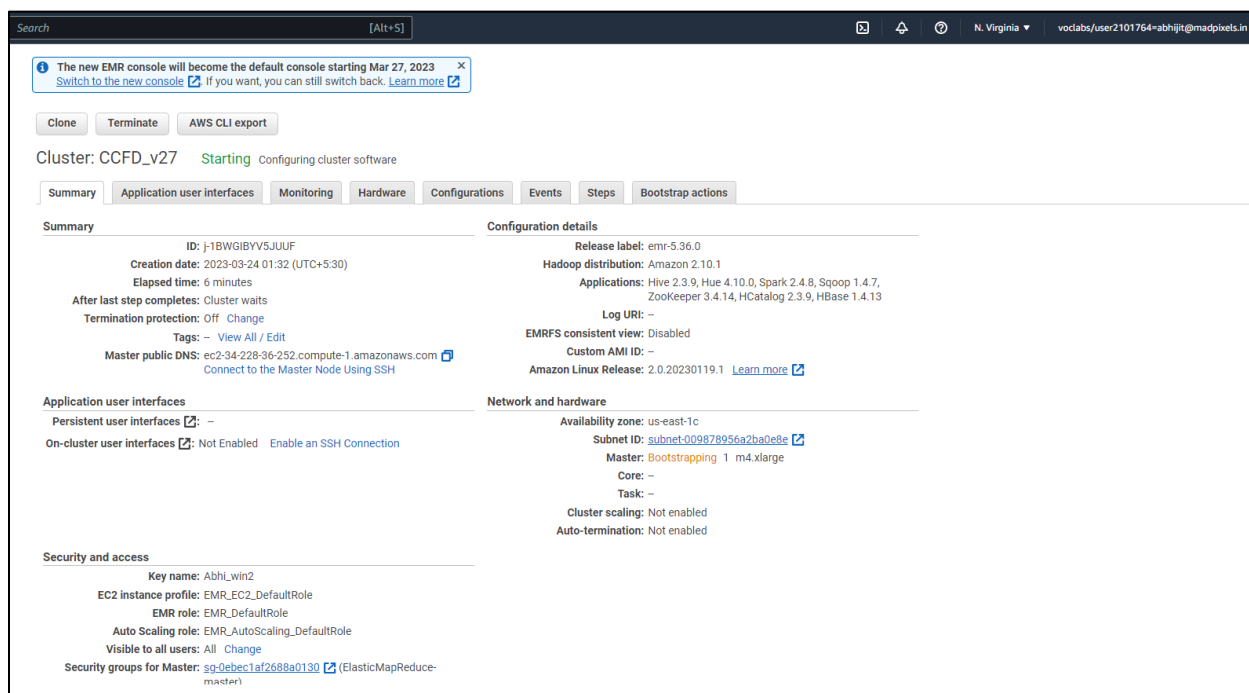


## Scripts Execution

### Explanation of the solution to the streaming layer problem

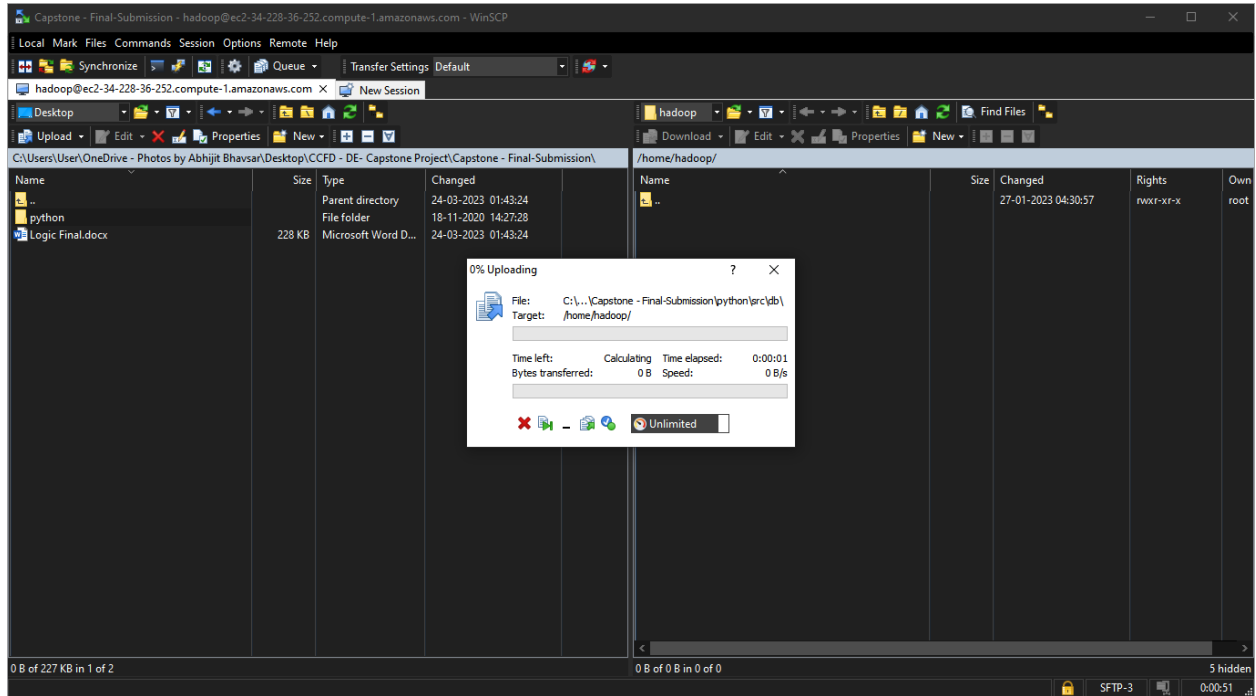
1. Created cluster.




The screenshot displays the AWS EMR console interface for a cluster named **CCFD\_v27**. The cluster is in a **Starting** state, with the status 'Configuring cluster software'. The console shows various tabs for managing the cluster, including Summary, Application user interfaces, Monitoring, Hardware, Configurations, Events, Steps, and Bootstrap actions. The Summary tab is selected, showing details such as the ID (j-1BWGIBYV5JUUF), Creation date (2023-03-24 01:32 UTC+5:30), Elapsed time (6 minutes), and the Master public DNS (ec2-34-228-36-252.compute-1.amazonaws.com). The Configuration details section shows the Release label (emr-5.36.0), Hadoop distribution (Amazon 2.10.1), and Applications (Hive 2.3.9, Hue 4.10.0, Spark 2.4.8, Sqoop 1.4.7, ZooKeeper 3.4.14, HCatalog 2.3.9, HBase 1.4.13). The Network and hardware section shows the Availability zone (us-east-1c), Subnet ID (subnet-009878956a2ba0e8e), and Master (Bootstrapping 1 m4.xlarge). The Security and access section shows the Key name (Abhi\_win2), EC2 instance profile (EMR\_EC2\_DefaultRole), EMR role (EMR\_DefaultRole), Auto Scaling role (EMR\_AutoScaling\_DefaultRole), and Visible to all users (All).

*Console Screenshot 1: Created EMR with Hadoop, Hive, HBase, HCatalog, Spark, ZooKeeper, Hue and Sqoop.*

## 2. Transferring Sample data to the cluster using WinSCP.



Console Screenshot 2: Copying data from local machine.



Console Screenshot 3: Copied the sample data provided by Upgrad.

3. After importing data into Hadoop Cluster, Editing **Dao.py** file

[illegible]

Console Screenshot 4: Editing dao.py file.

- Replacing self.host with Public IP address of our Master node.

```
hadoop@ip-172-31-20-140:~/python/src/db
if HBaseDao.__instance == None:
    HBaseDao()
return HBaseDao.__instance

def __init__(self):
    if HBaseDao.__instance != None:
        raise Exception("This class is a singleton!")
    else:
        HBaseDao.__instance = self
        self.host = '34.228.36.252'
        for i in range(2):
            try:
                self.pool = happybase.ConnectionPool(size=3, host=self.host,
port=9090)
                break
            except:
                print("Exception in connecting HBase")

def get_data(self, key, table):
    for i in range(2):
        try:
            with self.pool.connection() as connection:
-- INSERT --
```

Console Screenshot 5: Replacing self-host IP address with master node public IP.

- Updating rules.py with:

```
lookup_table = 'lookup_data_hbase'
master_table = 'card_transactions_hbase'
```

```
from db.dao import HBaseDao
from db.geo_map import GEO_Map
from datetime import datetime
import uuid

# Create UDF functions
lookup_table = 'lookup_data_hbase'
master_table = 'card_transactions_hbase'
speed_threshold = 0.25 # km/sec - Average speed of flight 900 km/hr
```

Console Screenshot 6: Created UDF function.

- Created Python functions, containing the logic for the UDFs (rules.py)  
verify\_ucl\_data : Function to verify the UCL rule Transaction amount should be less than Upper control limit (UCL)

```
def verify_ucl_data(card_id, amount):  
    try:  
        hbasedao = HBaseDao.get_instance()  
  
        card_row = hbasedao.get_data(key=str(card_id), table=lookup_table)  
        card_ucl = (card_row[b'card_data:ucl']).decode("utf-8")  
  
        if amount < float(card_ucl):  
            return True  
        else:  
            return False  
    except Exception as e:  
        raise Exception(e)
```

- verify\_credit\_score\_data: Function to verify the credit score rule .Credit score of each member should be greater than 200

```
def verify_credit_score_data(card_id):  
    try:  
        hbasedao = HBaseDao.get_instance()  
  
        card_row = hbasedao.get_data(key=str(card_id), table=lookup_table)  
        card_score = (card_row[b'card_data:score']).decode("utf-8")  
  
        if int(card_score) > 200:  
            return True  
        else:  
            return False  
    except Exception as e:  
        raise Exception(e)
```

8. verify\_postcode\_data: Function to verify the following zipcode rules. ZIP code distance.

```
hadoop@ip-172-31-20-140:~/python/src/rules
def verify_postcode_data(card_id, postcode, transaction_dt):
    try:
        hbasedao = HBaseDao.get_instance()
        geo_map = GEO_Map.get_instance()

        card_row = hbasedao.get_data(key=str(card_id), table=lookup_table)
        last_postcode = (card_row[b'card_data:postcode']).decode("utf-8")
        last_transaction_dt = (card_row[b'card_data:transaction_dt']).decode("utf-8")

        current_lat = geo_map.get_lat(str(postcode))
        current_lon = geo_map.get_long(str(postcode))
        previous_lat = geo_map.get_lat(last_postcode)
        previous_lon = geo_map.get_long(last_postcode)

        dist = geo_map.distance(lat1=current_lat, long1=current_lon, lat2=previous_lat, long2=previous_lon)

        speed = calculate_speed(dist, transaction_dt, last_transaction_dt)

        if speed < speed_threshold:
            return True
        else:
            return False

    except Exception as e:
        raise Exception(e)
```

9. calculate\_speed: A function to calculate the speed from distance and transaction timestamp differentials.

```
def calculate_speed(dist, transaction_dt1, transaction_dt2):

    transaction_dt1 = datetime.strptime(transaction_dt1, '%d-%m-%Y %H:%M:%S')
    transaction_dt2 = datetime.strptime(transaction_dt2, '%d-%m-%Y %H:%M:%S')

    elapsed_time = transaction_dt1 - transaction_dt2
    elapsed_time = elapsed_time.total_seconds()

    try:
        return dist / elapsed_time
    except ZeroDivisionError:
        return 299792.458
# (Speed of light)
```

10. `verify_rules_status`: A function to verify all the three rules - ucl, credit score and speed

```

hadoop@ip-172-31-20-140:~/python/src/rules
'''
def verify_rules_status(card_id, member_id, amount, pos_id, postcode, transaction_dt)
:

    hbasedao = HBaseDao.get_instance()

    # Check if the POS transaction passes all rules.
    # If yes, update the lookup table and insert data in master table as genuine.
    # Else insert the transaction in master table as Fraud.

    rule1 = verify_ucl_data(card_id, amount)
    rule2 = verify_credit_score_data(card_id)
    rule3 = verify_postcode_data(card_id, postcode, transaction_dt)

    if all([rule1, rule2, rule3]):
        status = 'GENUINE'
        hbasedao.write_data(key=str(card_id),
                            row={'card_data:postcode': str(postcode), 'card_data:tran
saction_dt': str(transaction_dt)},
                            table=lookup_table)

    else:
        status = 'FRAUD'

        new_id = str(uuid.uuid4()).replace('-', '')
        hbasedao.write_data(key=new_id,
                            row={'cardDetail:card_id': str(card_id), 'cardDetail:member_i
d': str(member_id),
                                'transactionDetail:amount': str(amount), 'transactionDet
ail:pos_id': str(pos_id),
                                'transactionDetail:postcode': str(postcode), 'transactio
nDetail:status': str(status),
                                'transactionDetail:transaction_dt': str(transaction_dt)}
                            ,
                            table=master_table)

    return status
'''

```

11. Now we are updating the **driver.py**

- Importing dependencies and setting Kafka consumer.
- Connecting to Kafka use the following details:  
**Bootstrap-server: 18.211.252.152**  
**Port Number: 9092**  
**Topic: transactions-topic-verified**
- Reading Input from Kafka

```
import os
import sys
from pyspark.sql import SparkSession
from pyspark.sql.functions import *
from pyspark.sql.types import *
from rules.rules import *

#initialising Spark session
spark = SparkSession \
    .builder \
    .appName("CreditCardFraud") \
    .getOrCreate()
spark.sparkContext.setLogLevel('ERROR')

# Reading input from Kafka
credit_data = spark.readStream \
    .format("kafka") \
    .option("kafka.bootstrap.servers", "18.211.252.152:9092") \
    .option("startingOffsets", "earliest") \
    .option("failOnDataLoss", "false") \
    .option("subscribe", "transactions-topic-verified") \
    .load()
```



12. Defining JSON schema of each transaction & reading the raw JSON data from Kafka as 'credit\_data\_stream' and Defining UDF's to verify rules.

```
# Defining schema for transaction
dataSchema = StructType() \
    .add("card_id", LongType()) \
    .add("member_id", LongType()) \
    .add("amount", DoubleType()) \
    .add("pos_id", LongType()) \
    .add("postcode", IntegerType()) \
    .add("transaction_dt", StringType())

# Casting raw data as string and aliasing
credit_data = credit_data.selectExpr("cast(value as string)")
credit_data_stream = credit_data.select(from_json(col="value", schema=dataSchema)
    .alias("credit_data")).select(
    "credit_data.*")

# Define UDF which verifies all the rules for each transaction and updates the lookup and master tables
verify_all_rules = udf(verify_rules_status, StringType())

Final_data = credit_data_stream \
    .withColumn('status', verify_all_rules(credit_data_stream['card_id'],
        credit_data_stream['member_id'],
        credit_data_stream['amount'],
        credit_data_stream['pos_id'],
        credit_data_stream['postcode'],
        credit_data_stream['transaction_dt']))

# Write output to console as well
```

13. Write output data to console.

```
# Write output to console as well
output_data = Final_data \
    .select("card_id", "member_id", "amount", "pos_id", "postcode", "transaction_dt") \
    .writeStream \
```

14. Spark Termination.

```
#indicating Spark to await termination
output_data.awaitTermination()
```

## 15. Installing Kafka Python using Root privileges

**Sudo -i**

**pip install kafka-python**

```

root@ip-172-31-20-140:~
EE::::EEEEEEEEEE::E M:::::M M:::::M R::::RRRRRR::::R
E::::E EEEEE M:::::M M:::::M RR::::R R::::R
E::::E M:::::M M:::::M M:::::M R::::R R::::R
E::::EEEEEEEEEE M::::M M::M M::M M::::M R::RRRRRR::::R
E:::::E M::::M M::M::M M::::M R:::::RR
E::::EEEEEEEEEE M::::M M::::M M::::M R::RRRRRR::::R
E::::E M::::M M::M M::::M R::R R::::R
E::::E EEEEE M::::M MMM M::::M R::R R::::R
EE::::EEEEEEEEEE::E M::::M M::::M R::R R::::R
E:::::E M::::M M::::M RR::::R R::::R
EEEEEEEEEEEEEEEEEEEE MMMMMM MMMMMM RRRRRRR RRRRRR

[root@ip-172-31-20-140 ~]# pip install kafka-python
WARNING: Running pip install with root privileges is generally not a good id
ry `pip3 install --user` instead.
Collecting kafka-python
  Downloading kafka-python-2.0.2-py2.py3-none-any.whl (246 kB)
    | 246 kB 36.4 MB/s
Installing collected packages: kafka-python
Successfully installed kafka-python-2.0.2
[root@ip-172-31-20-140 ~]#

```

Console Screenshot 7: Installing Kafka python

## 16. Updating all the Linux packages

**sudo yum update -y**

## 17. Installing Happybase & Pandas

```
sudo yum install python3-devel -y
```

```
pip install happybase
```

```
pip install pandas
```

```
hadoop@ip-172-31-20-140:~/python/src
python3-rpm-macros.noarch 0:3-60.amzn2.0.1

Complete!
[hadoop@ip-172-31-20-140 src]$ pip install happybase
Defaulting to user installation because normal site-packages is not writeable
Collecting happybase
  Downloading happybase-1.2.0.tar.gz (40 kB)
    |████████████████████████████████████████| 40 kB 6.3 MB/s
Requirement already satisfied: six in /usr/local/lib/python3.7/site-packages
(from happybase) (1.13.0)
Collecting thriftpy2>=0.4
  Downloading thriftpy2-0.4.16.tar.gz (643 kB)
    |████████████████████████████████████████| 643 kB 48.9 MB/s
Collecting ply<4.0,>=3.4
  Downloading ply-3.11-py2.py3-none-any.whl (49 kB)
    |████████████████████████████████████████| 49 kB 10.6 MB/s
Using legacy 'setup.py install' for happybase, since package 'wheel' is not
installed.
Using legacy 'setup.py install' for thriftpy2, since package 'wheel' is not
installed.
```

18. Making sure the permissions to directory for thrift services and starting the thrift service.

```
ls -ld /usr/lib/hbase/bin/./logs/
```

```
sudo chmod 777 /usr/lib/hbase/bin/./logs/
```

```
sudo touch /usr/lib/hbase/bin/./logs/hbase-hadoop-thrift-ip-172-31-20-140.out
```

```
/usr/lib/hbase/bin/hbase-daemon.sh start thrift -p 9090
```

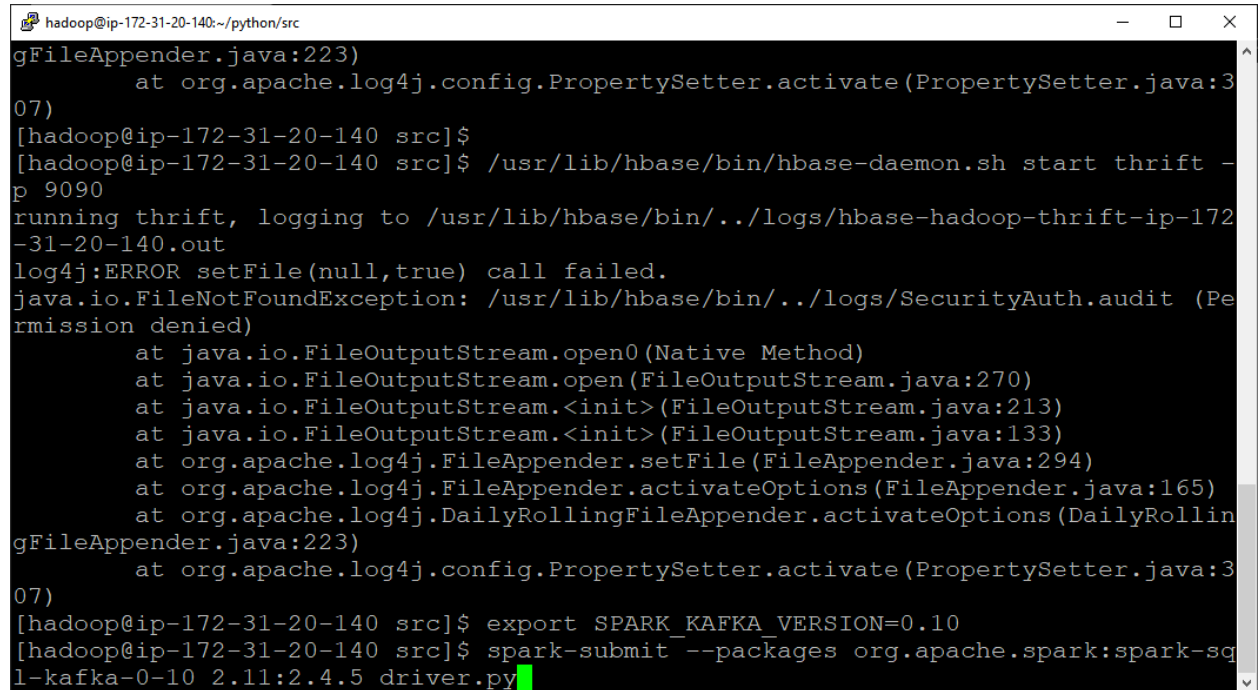
```
hadoop@ip-172-31-20-140:~/python/src
E:::E M:::M M:::M M:::M R:::R R:::R
E:::E EEEEE M:::M MMM M:::M R:::R R:::R
EE:::EEEEEE:::E M:::M M:::M R:::R R:::R
E:::E M:::M M:::M RR:::R R:::R
EEEEEEEEEEEEEEEE MMMMMM MMMMMM RRRRRR RRRRRR

[hadoop@ip-172-31-20-140 ~]$ cd python/src
[hadoop@ip-172-31-20-140 src]$ pip install happybase
Defaulting to user installation because normal site-packages is not writeable
Requirement already satisfied: happybase in /home/hadoop/.local/lib/python3.7/site-packages (1.2.0)
Requirement already satisfied: six in /usr/local/lib/python3.7/site-packages (from happybase) (1.13.0)
Requirement already satisfied: thriftpy2>=0.4 in /home/hadoop/.local/lib/python3.7/site-packages (from happybase) (0.4.16)
Requirement already satisfied: ply<4.0,>=3.4 in /home/hadoop/.local/lib/python3.7/site-packages (from thriftpy2>=0.4->happybase) (3.11)
[hadoop@ip-172-31-20-140 src]$ ls -ld /usr/lib/hbase/bin/./logs/
drwxr-xr-x 2 hbase hbase 4096 Mar 23 22:00 /usr/lib/hbase/bin/./logs/
[hadoop@ip-172-31-20-140 src]$ sudo chmod 777 /usr/lib/hbase/bin/./logs/
[hadoop@ip-172-31-20-140 src]$ sudo touch /usr/lib/hbase/bin/./logs/hbase-hadoop-thrift-ip-172-31-20-140.out
[hadoop@ip-172-31-20-140 src]$ /usr/lib/hbase/bin/hbase-daemon.sh start thrift -p 9090
```

**19.** Setting Kafka version and running spark submit command.

```
export SPARK_KAFKA_VERSION=0.10
```

```
spark-submit --packages org.apache.spark:spark-sql-kafka-0-10_2.11:2.4.5 driver.py
```

A terminal window screenshot showing the execution of a Spark submit command. The terminal title is 'hadoop@ip-172-31-20-140:~/python/src'. The output shows a Java exception: 'java.io.FileNotFoundException: /usr/lib/hbase/bin/../logs/SecurityAuth.audit (Permission denied)'. The user then runs 'export SPARK\_KAFKA\_VERSION=0.10' and 'spark-submit --packages org.apache.spark:spark-sql-kafka-0-10\_2.11:2.4.5 driver.py'.

```
hadoop@ip-172-31-20-140:~/python/src
gFileAppender.java:223)
    at org.apache.log4j.config.PropertySetter.activate(PropertySetter.java:307)
[hadoop@ip-172-31-20-140 src]$
[hadoop@ip-172-31-20-140 src]$ /usr/lib/hbase/bin/hbase-daemon.sh start thrift -p 9090
running thrift, logging to /usr/lib/hbase/bin/../logs/hbase-hadoop-thrift-ip-172-31-20-140.out
log4j:ERROR setFile(null,true) call failed.
java.io.FileNotFoundException: /usr/lib/hbase/bin/../logs/SecurityAuth.audit (Permission denied)
    at java.io.FileOutputStream.open0(Native Method)
    at java.io.FileOutputStream.open(FileOutputStream.java:270)
    at java.io.FileOutputStream.<init>(FileOutputStream.java:213)
    at java.io.FileOutputStream.<init>(FileOutputStream.java:133)
    at org.apache.log4j.FileAppender.setFile(FileAppender.java:294)
    at org.apache.log4j.FileAppender.activateOptions(FileAppender.java:165)
    at org.apache.log4j.DailyRollingFileAppender.activateOptions(DailyRollingFileAppender.java:133)
gFileAppender.java:223)
    at org.apache.log4j.config.PropertySetter.activate(PropertySetter.java:307)
[hadoop@ip-172-31-20-140 src]$ export SPARK_KAFKA_VERSION=0.10
[hadoop@ip-172-31-20-140 src]$ spark-submit --packages org.apache.spark:spark-sql-kafka-0-10_2.11:2.4.5 driver.py
```

Console Screenshot 8: Spark Submit Command

## 18. Count Data in Hbase: count 'lookup\_data\_hive'

```
hadoop@ip-172-31-20-140:~/python/src
request executors before the AM has registered!
23/03/23 22:23:26 INFO YarnClientSchedulerBackend: SchedulerBackend is ready for
scheduling beginning after reached minRegisteredResourcesRatio: 0.0
23/03/23 22:23:26 INFO YarnSchedulerBackend$YarnSchedulerEndpoint: ApplicationMa
ster registered as NettyRpcEndpointRef(spark-client://YarnAM)
23/03/23 22:23:26 INFO SharedState: loading hive config file: file:/etc/spark/co
nf.dist/hive-site.xml
23/03/23 22:23:26 INFO SharedState: Setting hive.metastore.warehouse.dir ('null'
) to the value of spark.sql.warehouse.dir ('hdfs:///user/spark/warehouse').
23/03/23 22:23:26 INFO SharedState: Warehouse path is 'hdfs:///user/spark/wareho
use'.
23/03/23 22:23:26 INFO JettyUtils: Adding filter org.apache.hadoop.yarn.server.w
ebproxy.amfilter.AmIpFilter to /SQL.
23/03/23 22:23:26 INFO JettyUtils: Adding filter org.apache.hadoop.yarn.server.w
ebproxy.amfilter.AmIpFilter to /SQL/json.
23/03/23 22:23:26 INFO JettyUtils: Adding filter org.apache.hadoop.yarn.server.w
ebproxy.amfilter.AmIpFilter to /SQL/execution.
23/03/23 22:23:26 INFO JettyUtils: Adding filter org.apache.hadoop.yarn.server.w
ebproxy.amfilter.AmIpFilter to /SQL/execution/json.
23/03/23 22:23:26 INFO JettyUtils: Adding filter org.apache.hadoop.yarn.server.w
ebproxy.amfilter.AmIpFilter to /static/sql.
23/03/23 22:23:27 INFO StateStoreCoordinatorRef: Registered StateStoreCoordinator
 endpoint
```

### Please Note:

1. Due to the issue with Kafka server, I could not be able to get any data. However, once Batch 0 has data printed into console. We should be able to verify the number of entries.
2. We also tried doing it with the static data which was shared by upgrad in text format. We bypassed Kafka and imported directly using spark streaming from the local storage but even after numerous attempts spark could not access or find the file. We checked the read and write permission as well, but we kept getting the error that file does not exist.

Code for which was replaced in with Kafka Consumer details:

```
# Reading input from a text file
```

```
credit_data = spark.readStream \
    .option("inferSchema", "true") \
    .text("/home/hadoop/python/src/transactions-topic-verified.txt")
```